VISUAL QUESTION ANSWETRING (VQA)

Foundations of Machine Learning – Fall 2018

Daniel Rivera Ruiz

New York University

Outline

1) What is VQA?
2) The VQA Dataset
3) LSTM I+Q
4) Hierarchical Co-Attention
5) Stacked Attention Networks
6) Multimodal Residual Learning
7) Focused Dynamic Attention
8) Dynamic Memory Networks
9) Multimodal Compact Bilinear
10) DualNet Architecture
11) Results Comparison
12) The VQA 2.0 Dataset
13) Future work

What is VQA?

Antol et al. (2015)

- Image + Question = Answer
- Intersection of Computer Vision and NLP
- "AI-complete" task:
  - ✓ multimodal knowledge
  - ✓ Evaluation metric

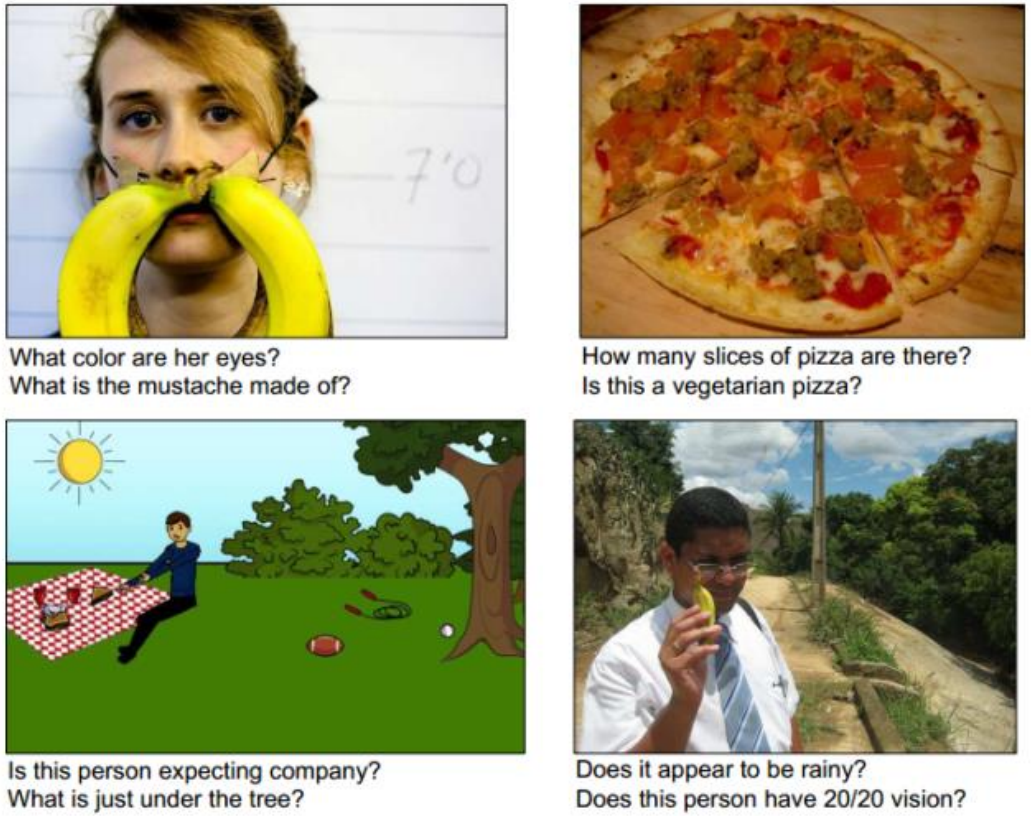

The problem

Antol et al. (2015)

Predict the most likely answer $\hat{a}$ for a given image **x** and a question or phrase **q**:

$$\hat{a} = \underset{a \in A}{\operatorname{argmax}}\ p(a|\boldsymbol{x}, \boldsymbol{q}; \theta)$$

with model parameters θ and a set of possible answers A

The dataset

Antol et al. (2015)

- 204,721 images from the MS COCO dataset
- 50,000 abstract scenes
- 3 questions per image (over 760,000)
- 10 answers per question
- 2 tasks: open-ended and multiple-choice
- Open-ended accuracy metric: $\min\left(\frac{\#\text{ humans that provided the answer}}{3}, 1\right)$

The dataset

Antol et al. (2015)

Plausible answers (blue) are collected for the multiple-choice task by having subjects answer the question without looking at the image

How many bikes are there?	2 2 2	3 4 12
What number is the bus?	48 48 48	4 46 number 6

The dataset

Antol et al. (2015)

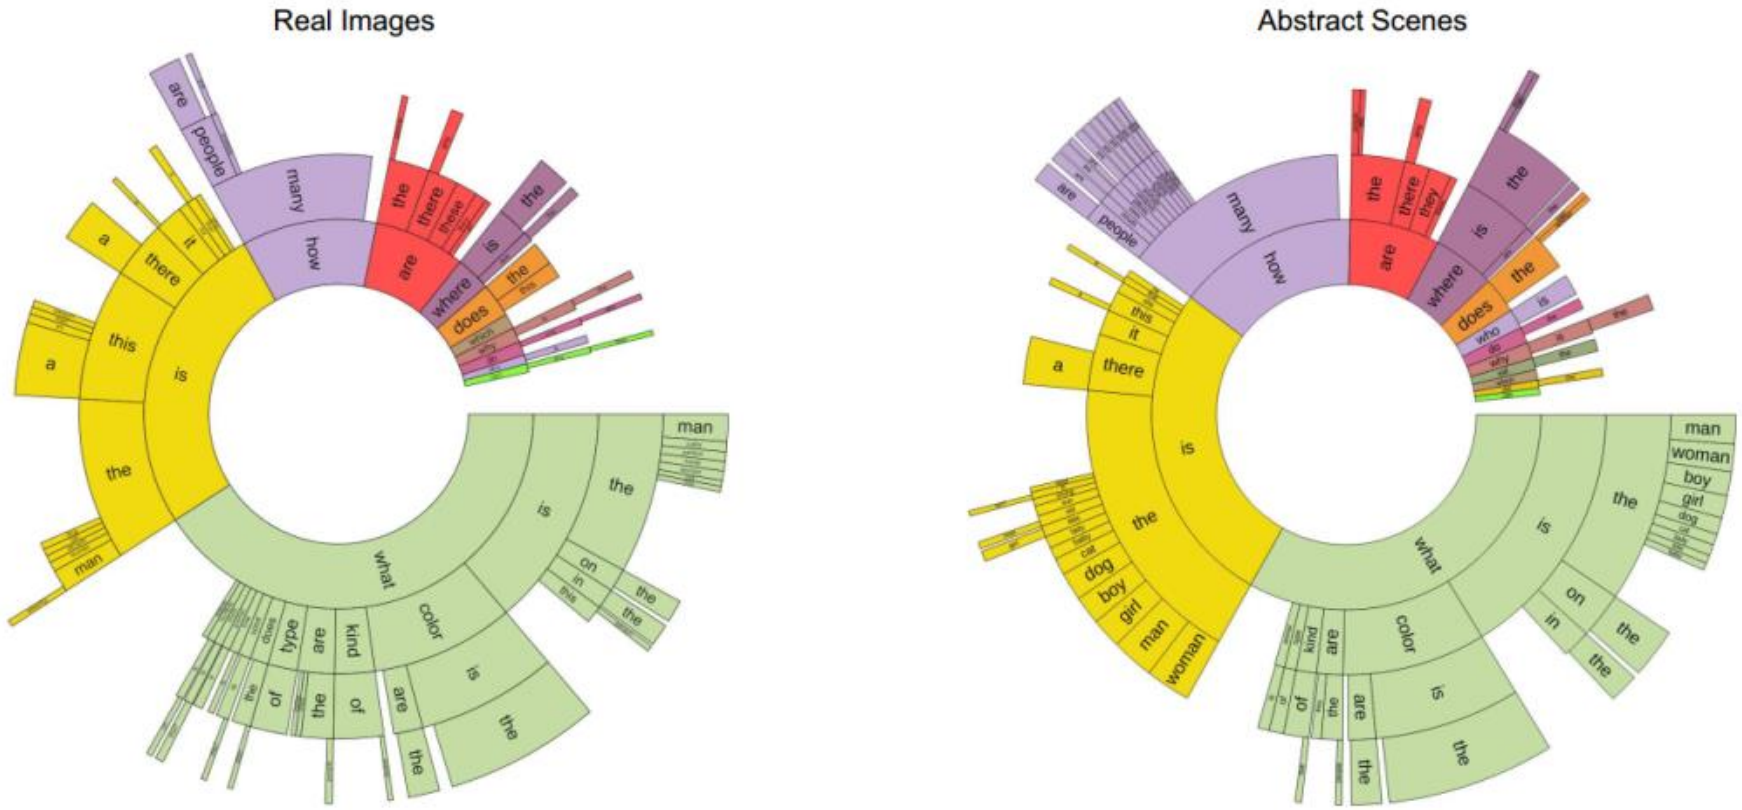

The dataset

Antol et al. (2015)

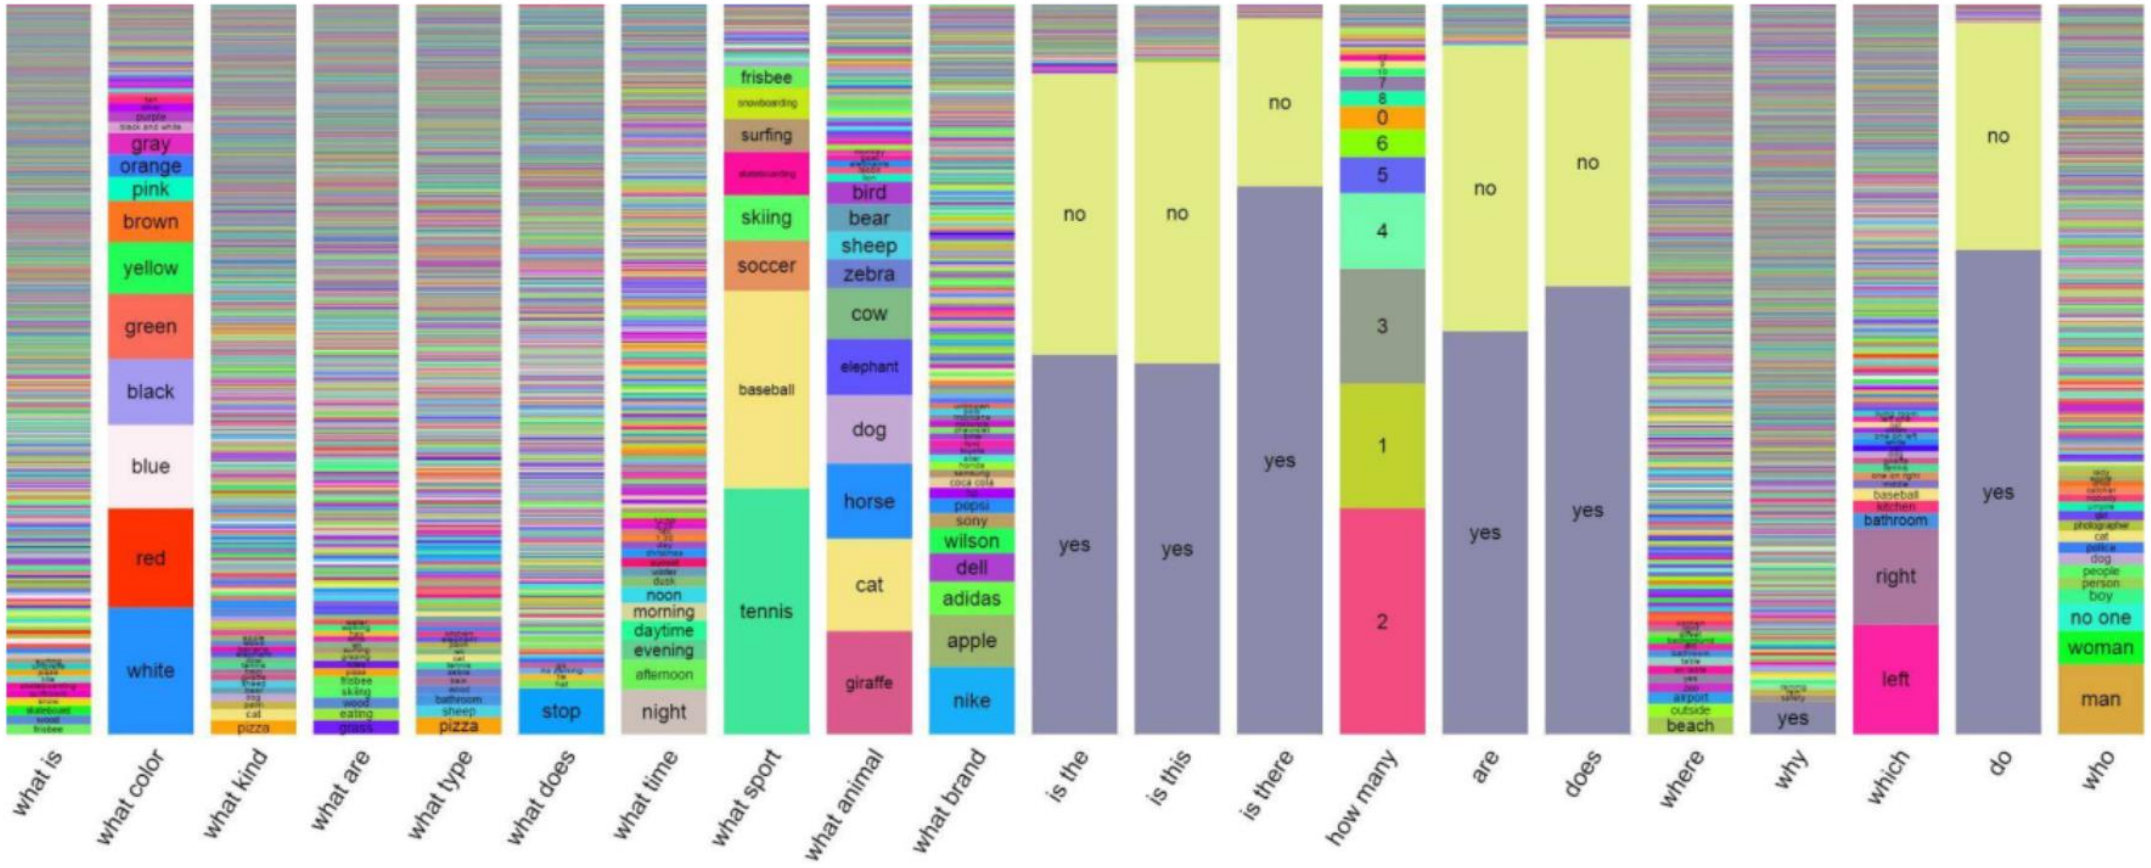

LSTM I+Q

Antol et al. (2015)

- Baseline model published with the VQA dataset
- Image features: last hidden layer (4096) of VGGNet
- Question features: LSTM (1024)
- Multimodality: element-wise multiplication (EWM)
- Attention mechanism: none
- Test-dev accuracy: **53.7%**
- Test-std accuracy: **54.1%**

Hierarchical Co-Attention

Lu et al. (2016)

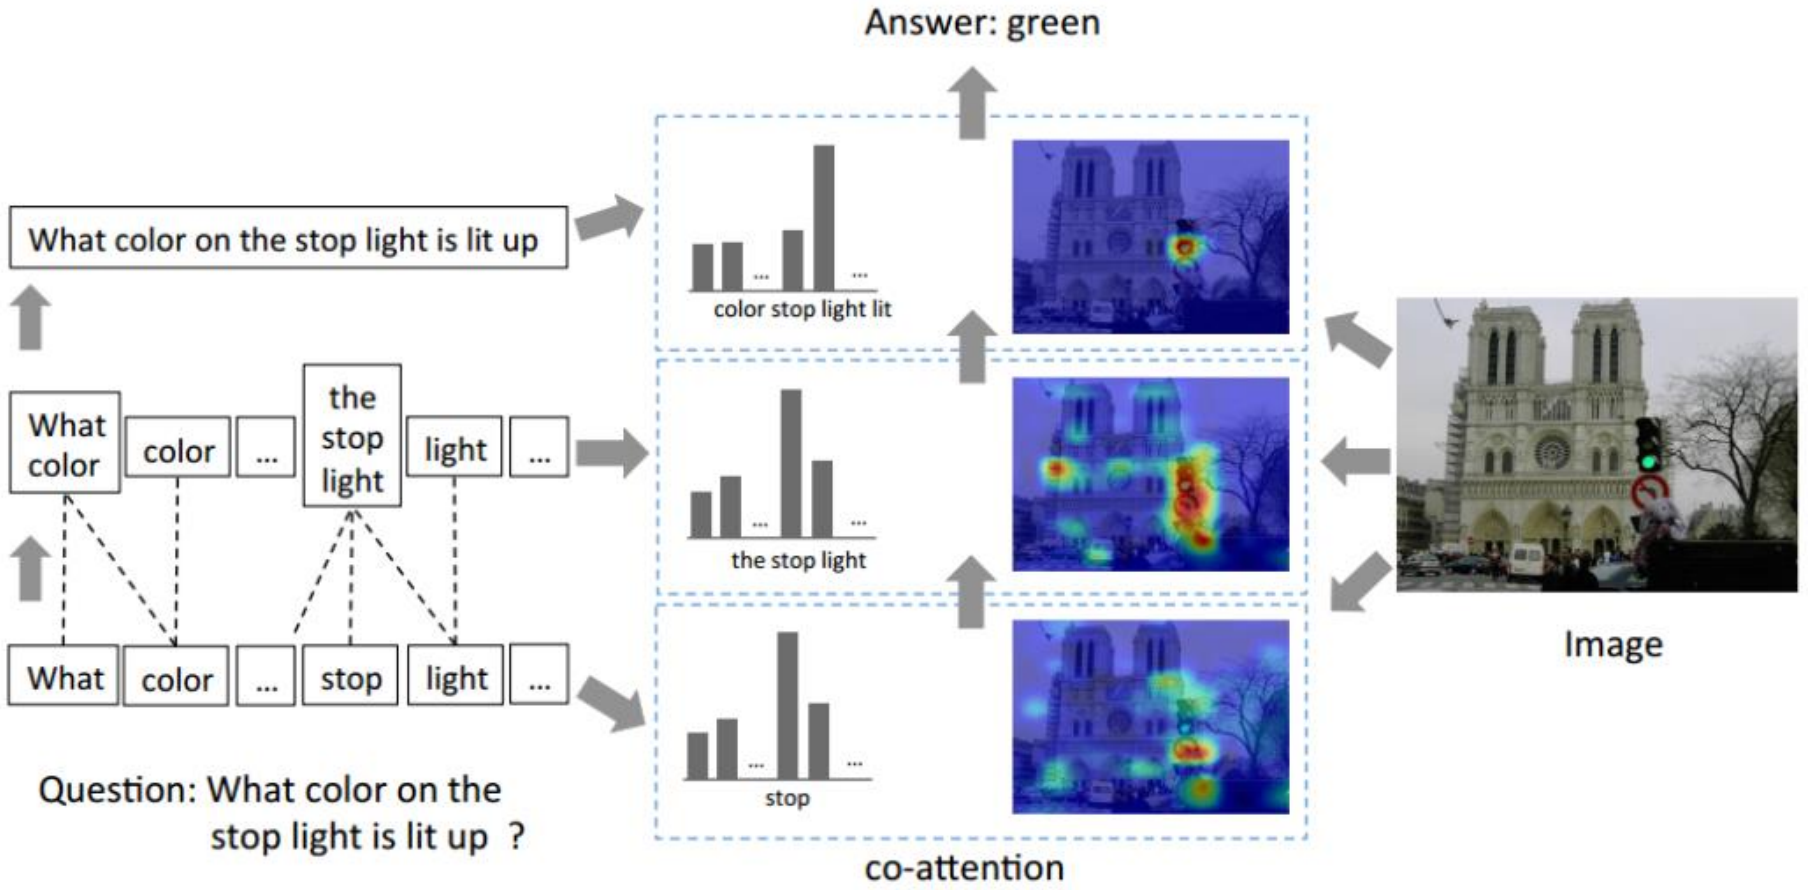

Hierarchical Co-Attention

Lu et al. (2016)

Question hierarchy

$$\boldsymbol{Q}^w = \{\boldsymbol{q}_1^w, \boldsymbol{q}_2^w, \ldots, \boldsymbol{q}_T^w\}$$

$$\hat{\boldsymbol{q}}_{s,t}^p = \tanh(\boldsymbol{W}_c^s \boldsymbol{q}_{t:t+s-1}^w),$$

$$\boldsymbol{q}_{s,t}^p = \max(\hat{\boldsymbol{q}}_{1,t}^p, \hat{\boldsymbol{q}}_{2,t}^p, \hat{\boldsymbol{q}}_{3,t}^p)$$

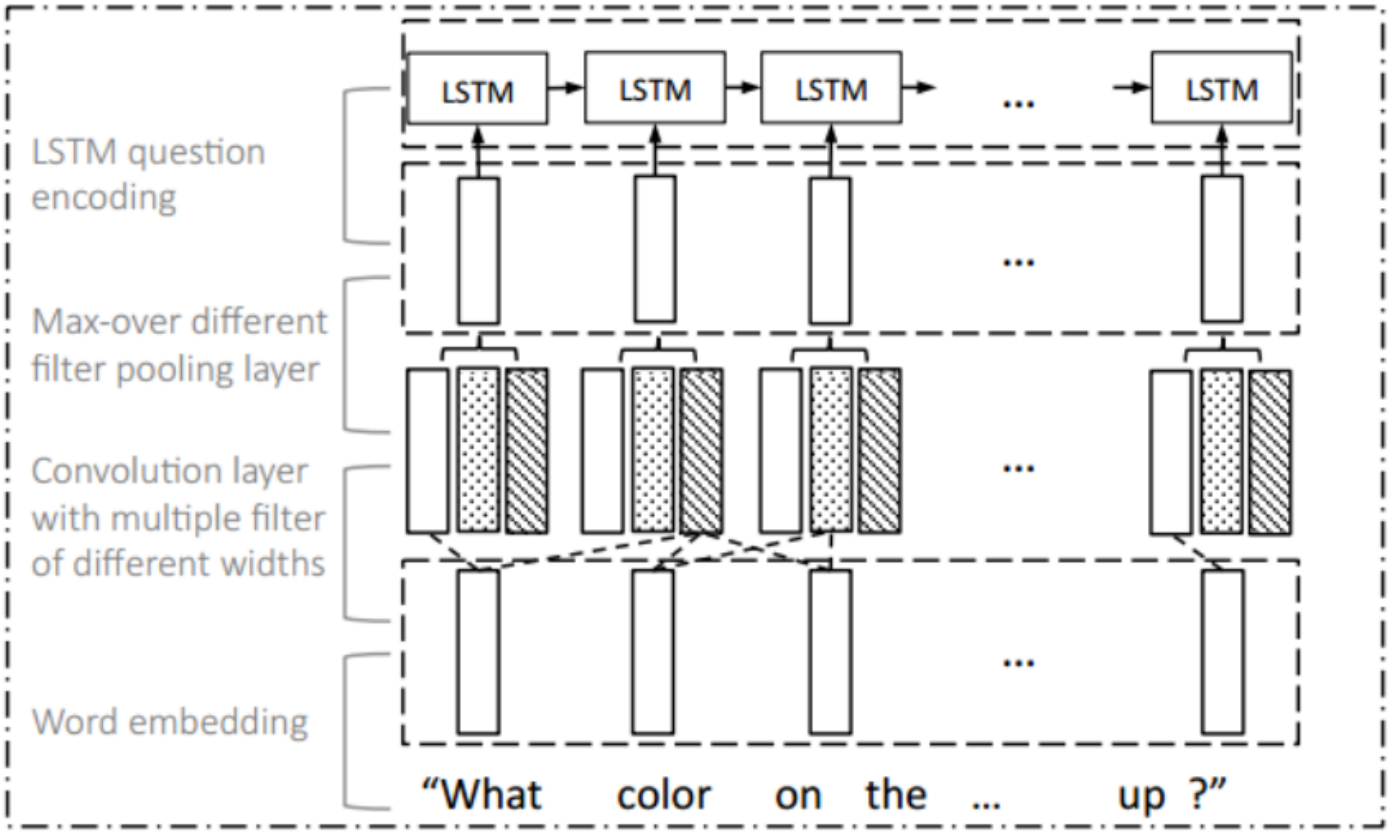

Hierarchical Co-Attention

Lu et al. (2016)

Parallel

$$\boldsymbol{C} = \tanh(\boldsymbol{Q}^T \boldsymbol{W}_b \boldsymbol{V})$$

$$\boldsymbol{H}^v = \tanh(\boldsymbol{W}_v \boldsymbol{V} + (\boldsymbol{W}_q \boldsymbol{Q})\boldsymbol{C})$$

$$\boldsymbol{a}^v = \text{softmax}(w_{hv}^T \boldsymbol{H}^v)$$

$$\boldsymbol{H}^q = \tanh(\boldsymbol{W}_q \boldsymbol{Q} + (\boldsymbol{W}_v \boldsymbol{V})\boldsymbol{C}^T)$$

$$\boldsymbol{a}^q = \text{softmax}(w_{hq}^T \boldsymbol{H}^q)$$

$$\hat{\boldsymbol{v}} = \sum_{n=1}^{N} a_n^v \boldsymbol{v}_n \qquad \hat{\boldsymbol{q}} = \sum_{t=1}^{T} a_t^q \boldsymbol{q}_t$$

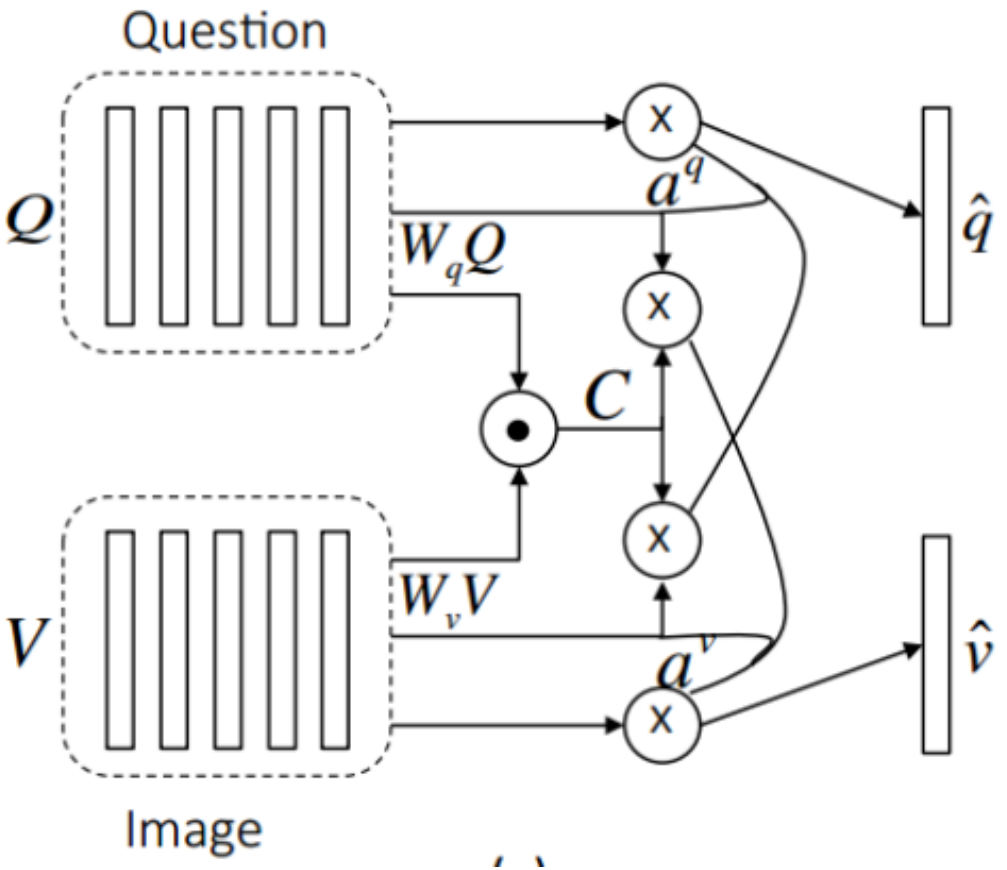

Hierarchical Co-Attention

Lu et al. (2016)

Alternating

$$\boldsymbol{H} = \tanh(\boldsymbol{W}_x \boldsymbol{X} + (\boldsymbol{W}_g \boldsymbol{g})\mathbf{1}^T)$$

$$\boldsymbol{a}^x = \text{softmax}(w_{hx}^T \boldsymbol{H})$$

$$\hat{\boldsymbol{x}} = \sum a_i^x \boldsymbol{x}_i$$

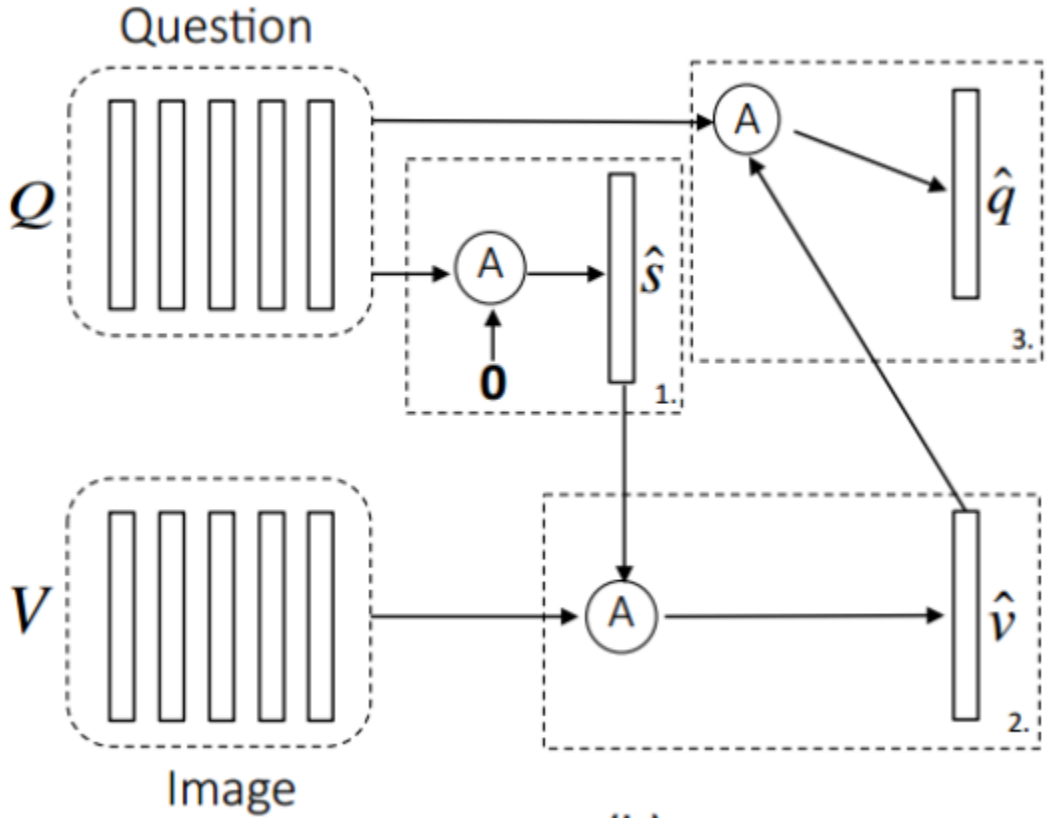

Hierarchical Co-Attention

Lu et al. (2016)

Prediction

$$
\begin{aligned}
\boldsymbol{h}^w &= \tanh(\boldsymbol{W}_w(\hat{\boldsymbol{q}}^w + \hat{\boldsymbol{v}}^w)) \\
\boldsymbol{h}^p &= \tanh(\boldsymbol{W}_p[(\hat{\boldsymbol{q}}^p + \hat{\boldsymbol{v}}^p); \boldsymbol{h}^w]) \\
\boldsymbol{h}^s &= \tanh(\boldsymbol{W}_s[(\hat{\boldsymbol{q}}^s + \hat{\boldsymbol{v}}^s); \boldsymbol{h}^p]) \\
p &= \text{softmax}(\boldsymbol{W}_h \boldsymbol{h}^s)
\end{aligned}
$$

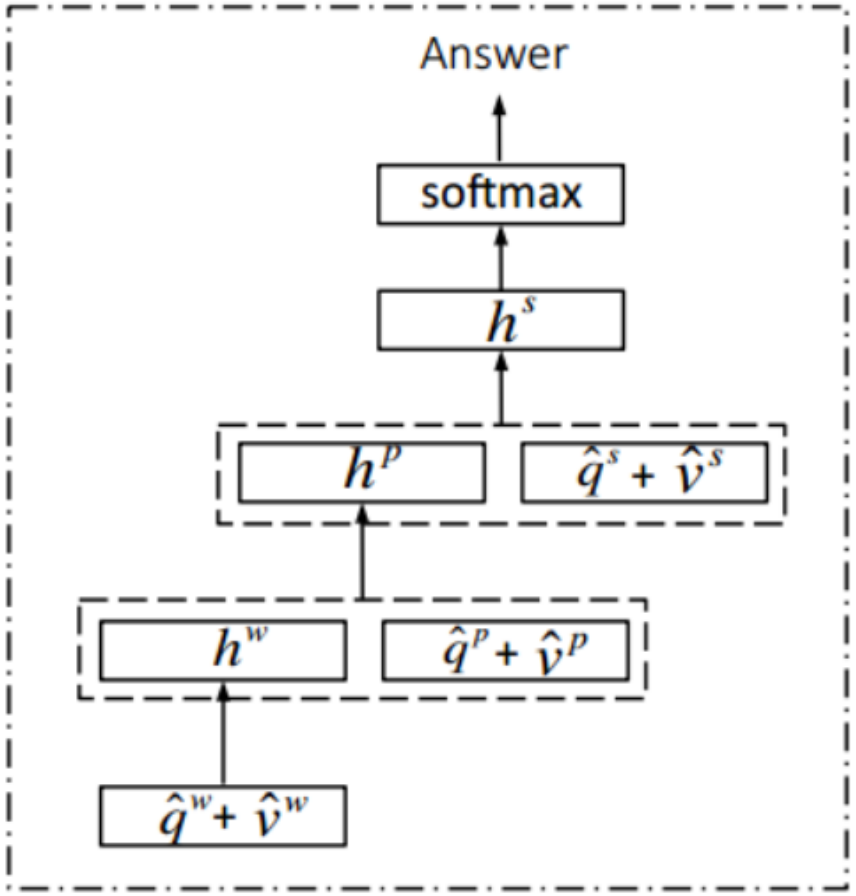

Hierarchical Co-Attention

Lu et al. (2016)

- Image features: last hidden layer (4096) of VGGNet
- Question features: LSTM (1024)
- Multimodality: EWM
- Attention mechanism: Image + Question
- Test-dev accuracy: **61.8%**
- Test-std accuracy: **62.1%**

Hierarchical Co-Attention

Lu et al. (2016)

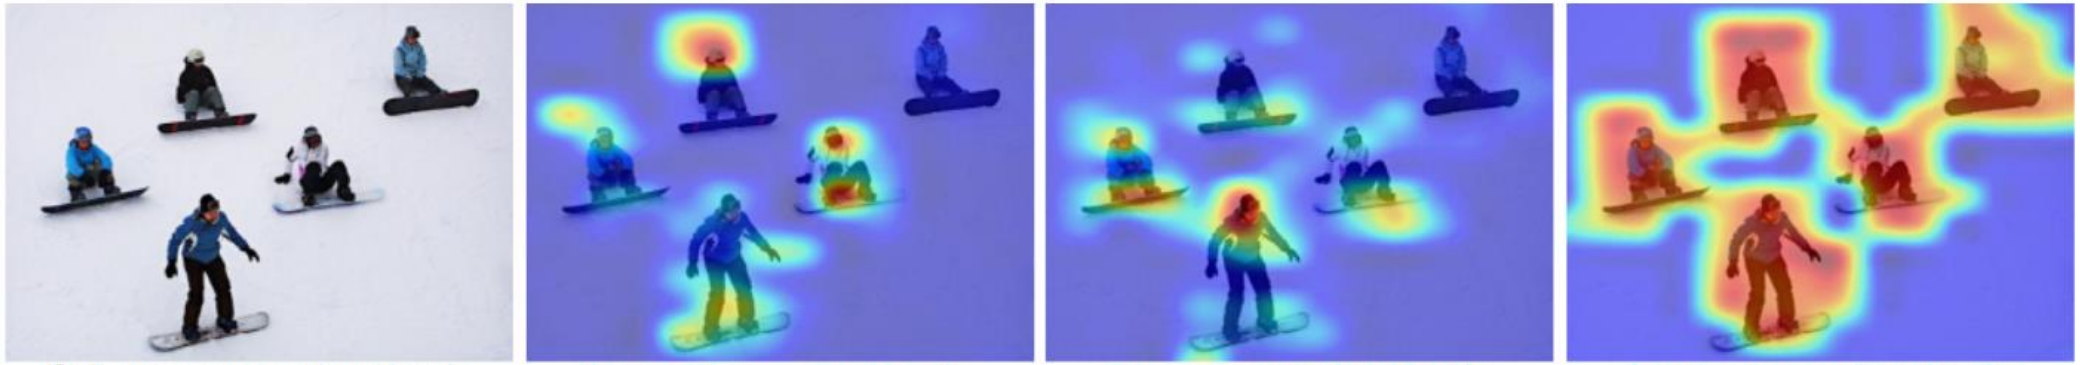

Stacked Attention Networks (SAN) Yang et al. (2015)

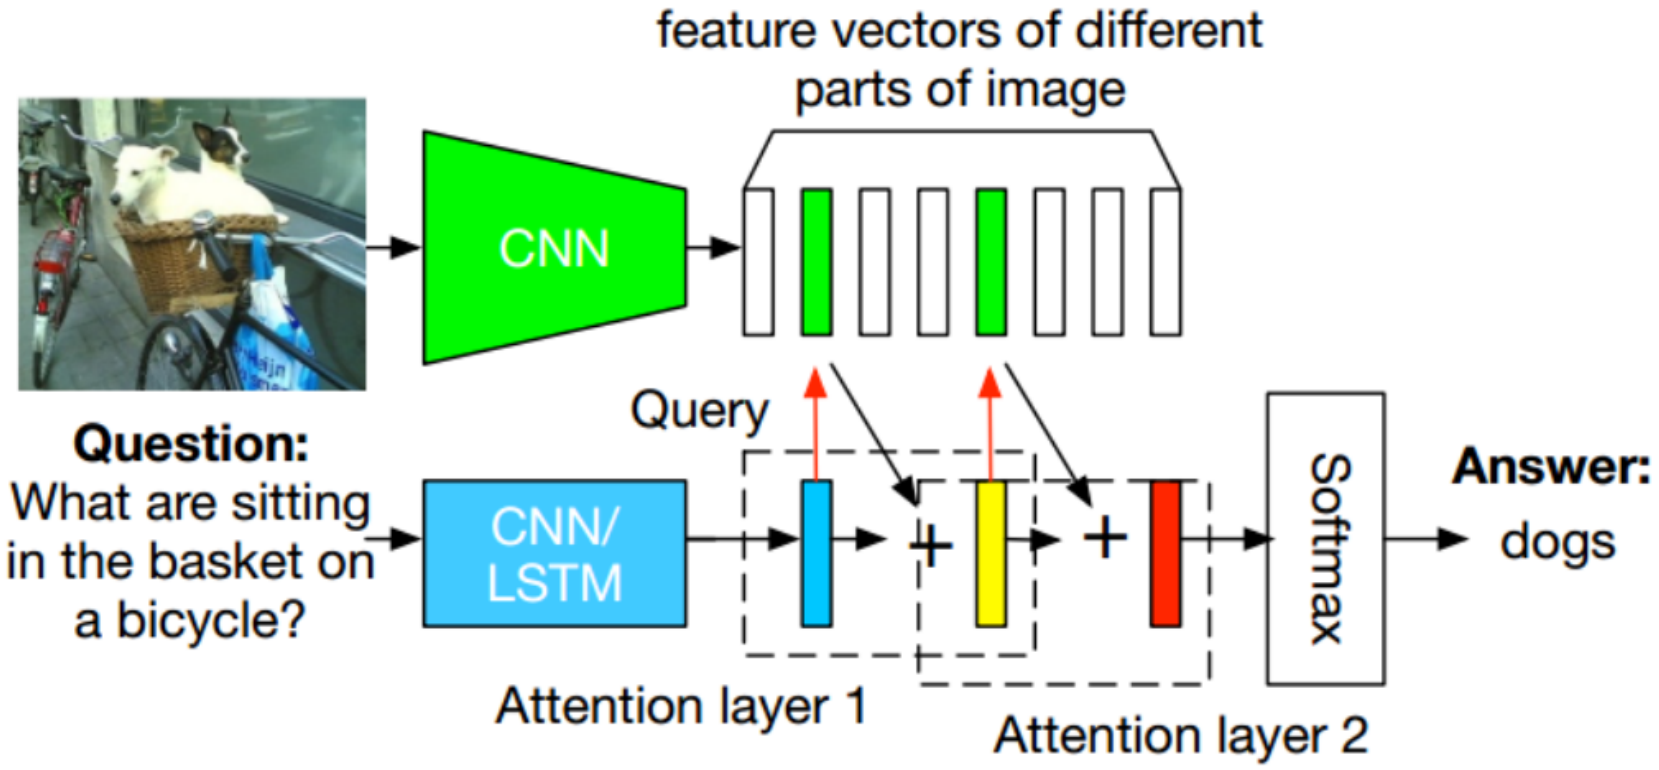

Stacked Attention Networks (SAN) Yang et al. (2015)

$$\tilde{h}_c = \max[h_{c,1}, h_{c,2}, \ldots, h_{c,T-c+1}]$$

$$h = [\tilde{h}_1, \tilde{h}_2, \tilde{h}_3]$$

$$h_A^k = \tanh(W_{I,A}^k v_I \oplus (W_{Q,A}^k u^{k-1} + b_A^k)$$

$$p_I^k = softmax(W_p^k h_A^k + b_p^k)$$

$$\tilde{v}_I^k = \sum p_i^k v_i$$

$$u^k = \tilde{v}_I^k + u^{k-1}$$

$$p_{ans} = \text{softmax}(W_u u^K + b_u)$$

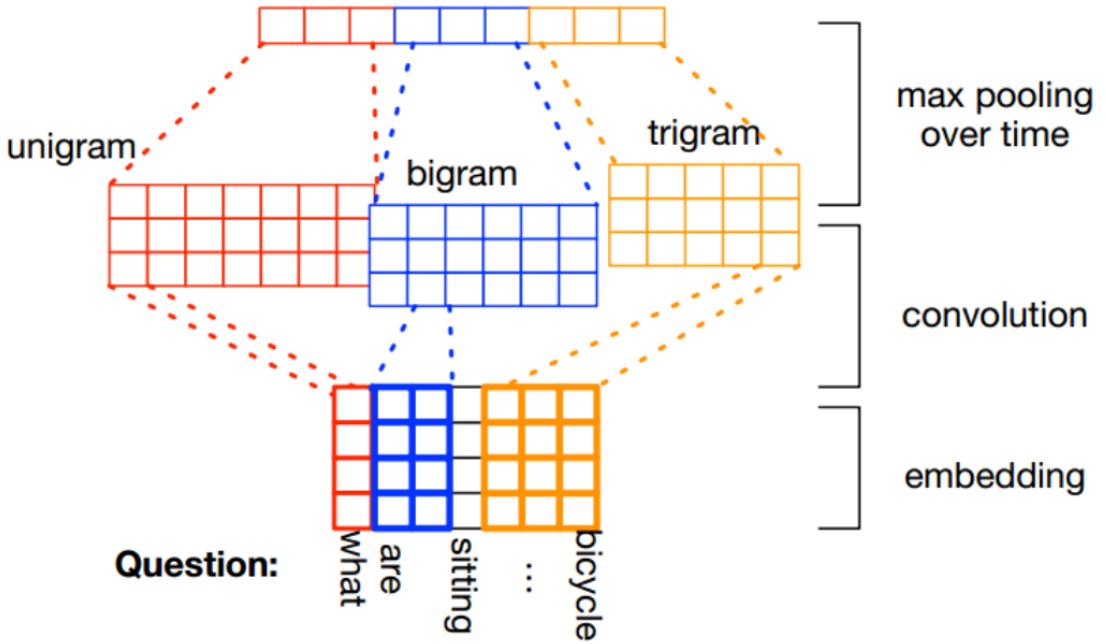

Stacked Attention Networks (SAN) Yang et al. (2015)

- Image features: VGGNet (512)
- Question features: CNN (1280)
- Multimodality: EWM
- Attention mechanism: Question
- Test-dev accuracy: **58.7%**
- Test-std accuracy: **58.9%**

Stacked Attention Networks (SAN) Yang et al. (2015)

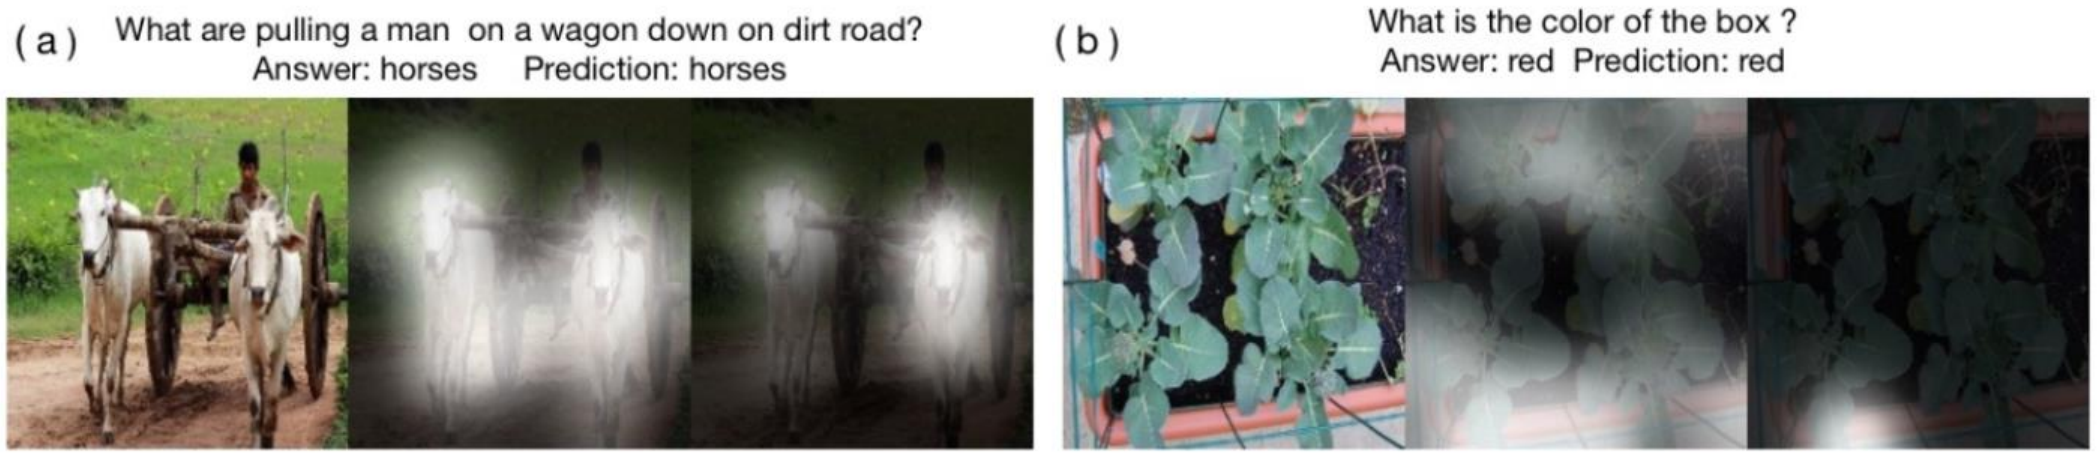

Multimodal Residual Network (MRN)

Kim et al. (2016)

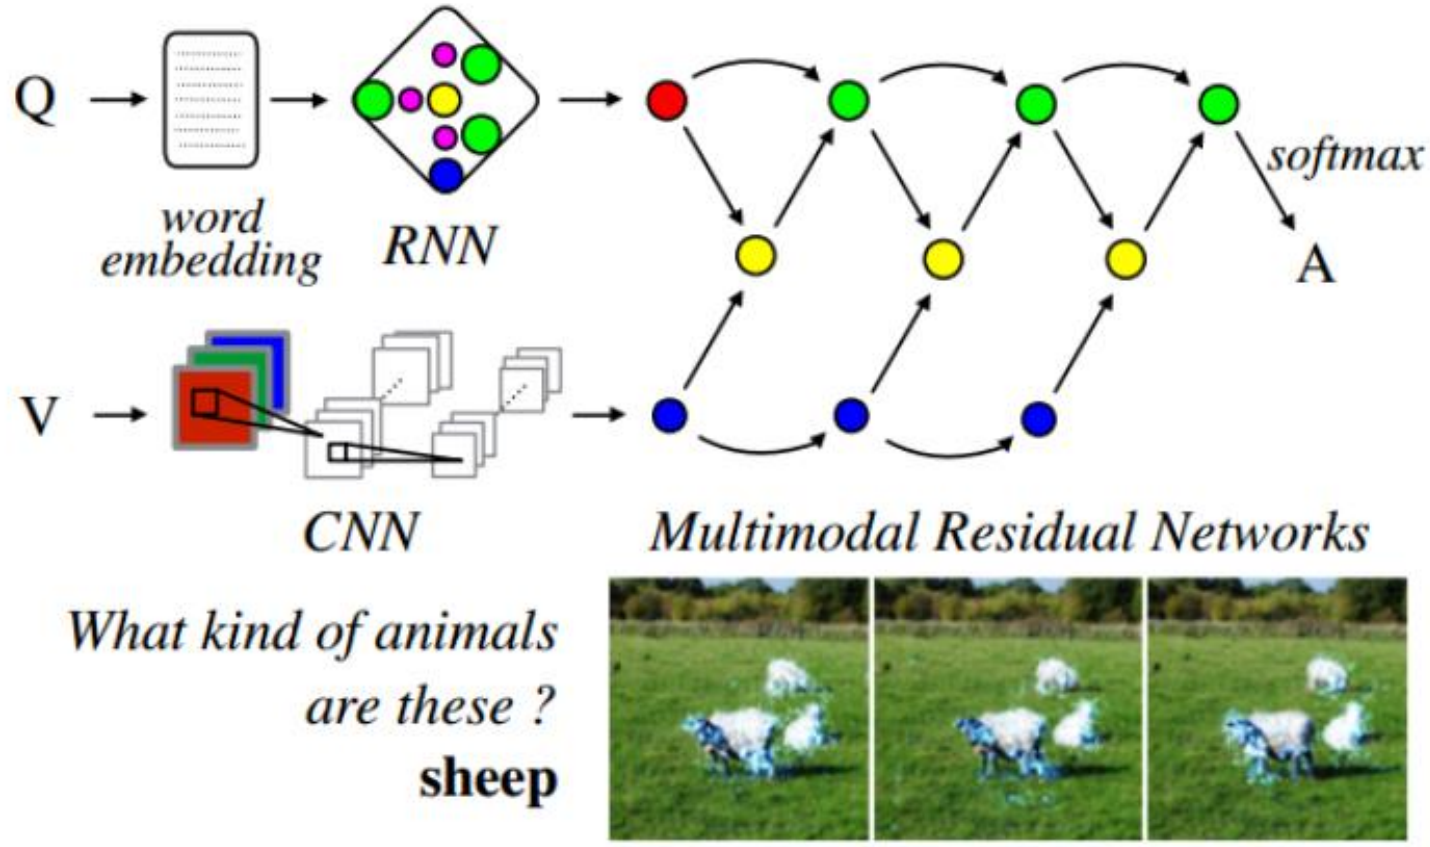

Multimodal Residual Network (MRN)

$$\boldsymbol{q}^k = F(\boldsymbol{q}^{k-1}, \boldsymbol{V}) + \boldsymbol{q}^{k-1}$$

$$F^{(k)}(\boldsymbol{q}, \boldsymbol{v}) = \tanh(W_{\boldsymbol{q}}^{(k)}\boldsymbol{q} \odot \tanh(W_2^{(k)}\tanh(W_1^{(k)}\boldsymbol{v}))$$

$$H_L(\boldsymbol{q}, \boldsymbol{v}) = W_{\boldsymbol{q}'}\boldsymbol{q} + \sum_{l=1}^{L} W_{F^{(l)}} F^{(l)}(H_{l-1}, \boldsymbol{v})$$

$$H_0 = \boldsymbol{q}$$

$$W_{\boldsymbol{q}'} = \prod_{l=1}^{L} W_{\boldsymbol{q}'}^{(l)} \qquad W_{F^{(l)}} = \prod_{m=l+1}^{L} W_{\boldsymbol{q}'}^{(m)}$$

Kim et al. (2016)

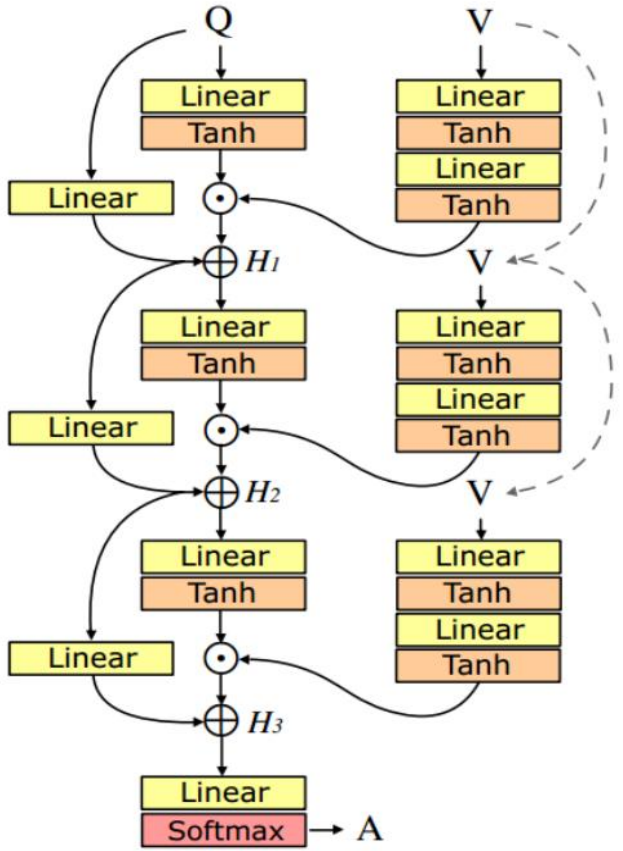

Multimodal Residual Network (MRN)

Kim et al. (2016)

- Image features: last hidden layer (4096) of VGGNet
- Question features: GRU (2400)
- Multimodality: EWM
- Attention mechanism: Implicit
- Test-dev accuracy: **60.5%**
- Test-std accuracy: **61.8%**

Multimodal Residual Network (MRN)

Kim et al. (2016)

(a) *What kind of animals are these ?* **sheep**

(b) *What animal is the picture ?* **elephant**

Focused Dynamic Attention (FDA)

Ilievski et al. (2016)

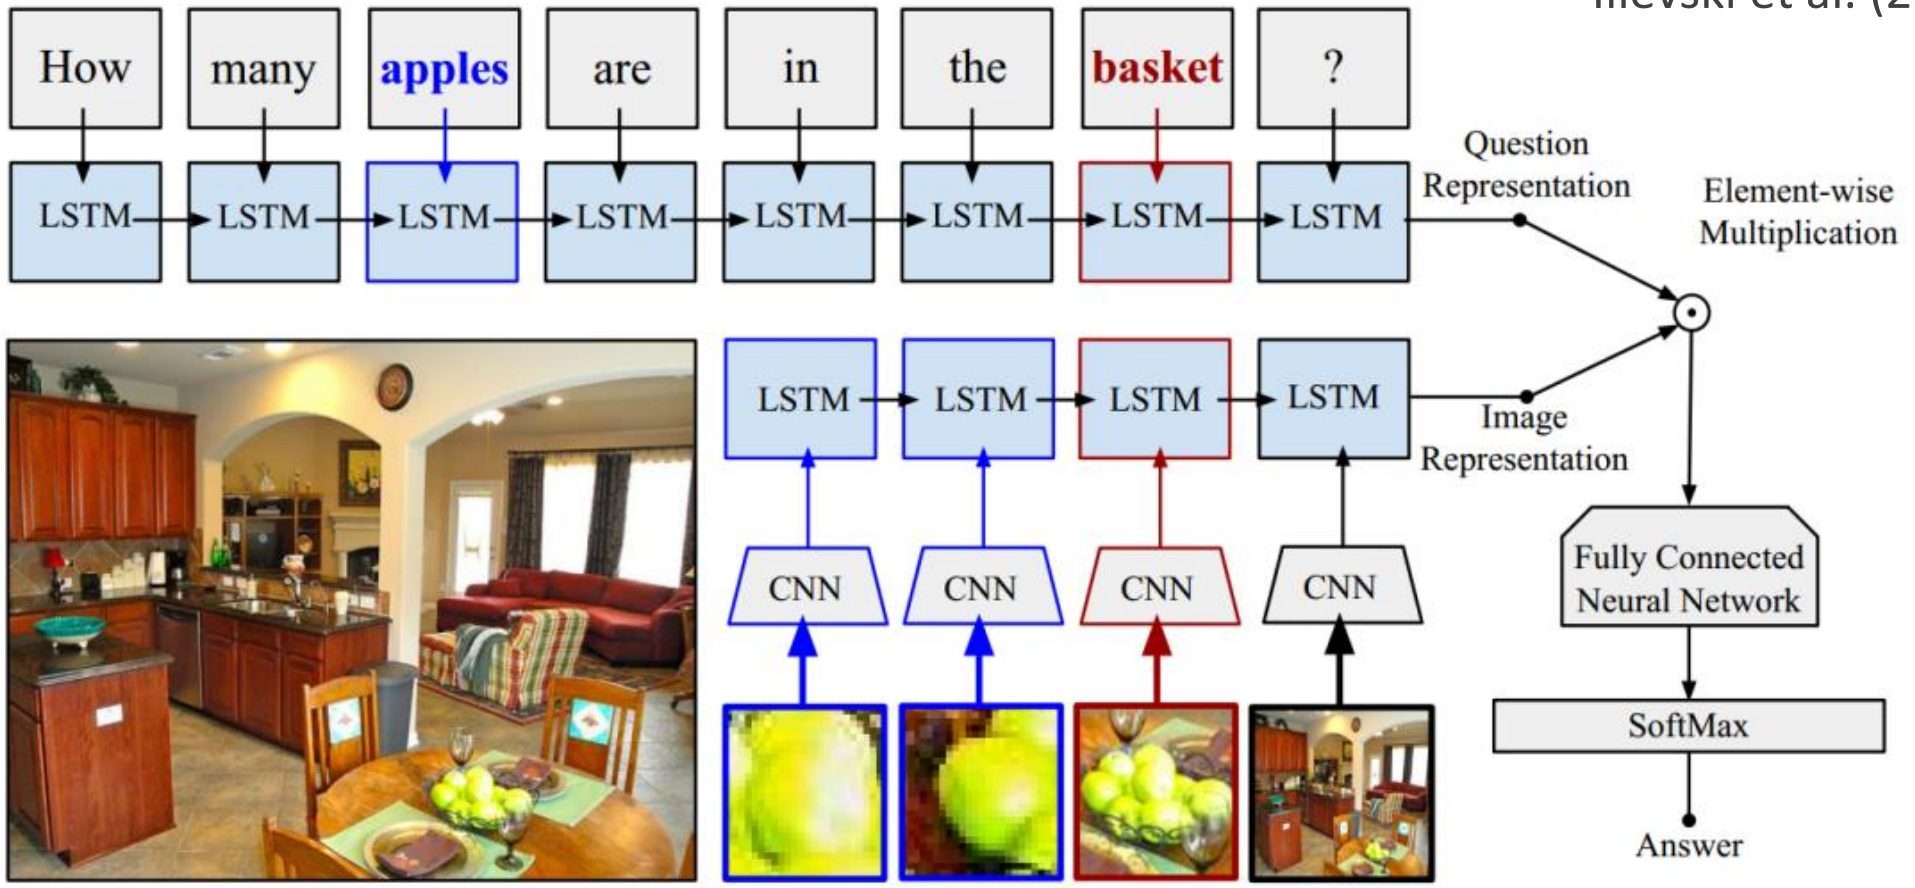

Focused Dynamic Attention (FDA)

Ilievski et al. (2016)

- Image features: ResNet (2048)
- Question features: LSTM (512)
- Multimodality: EWM
- Attention mechanism: Object detection
- Test-dev accuracy: **59.2%**
- Test-std accuracy: **59.5%**

Focused Dynamic Attention (FDA)

Ilievski et al. (2016)

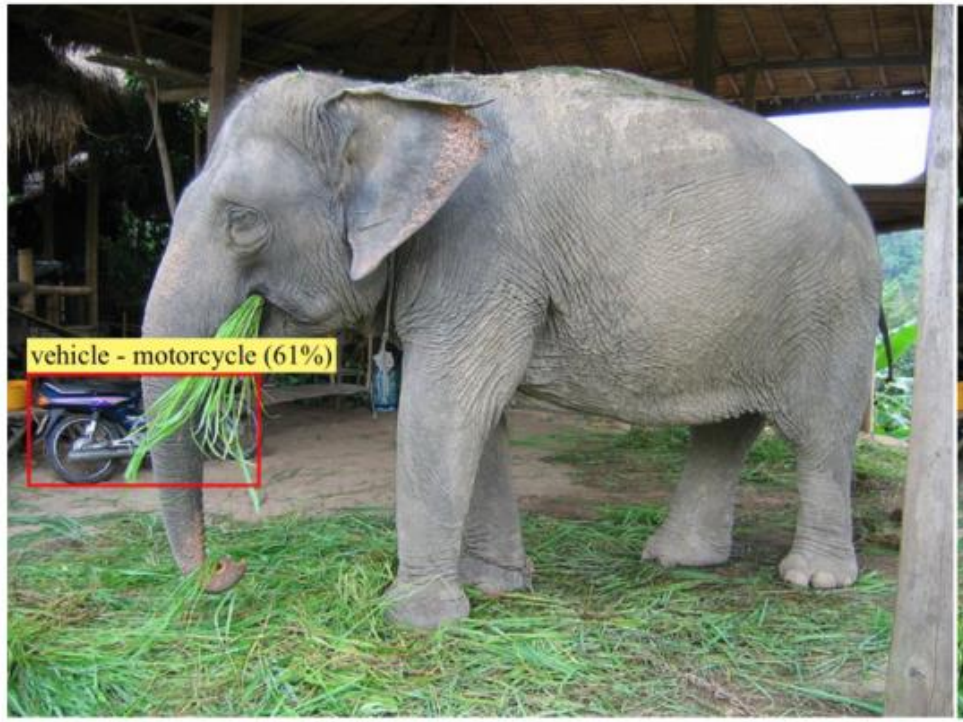

What type of **vehicle** is pictured?
- Motorcycle.

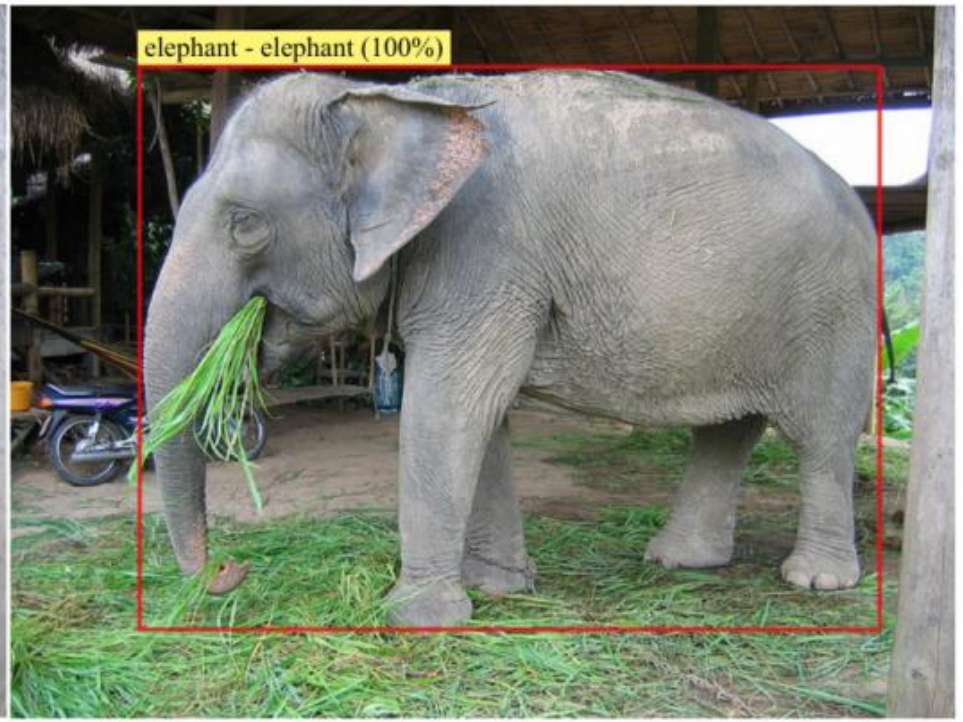

Does the **elephant** have tusks?
- No.

Dynamic Memory Networks (DMN)

Xiong et al. (2016)

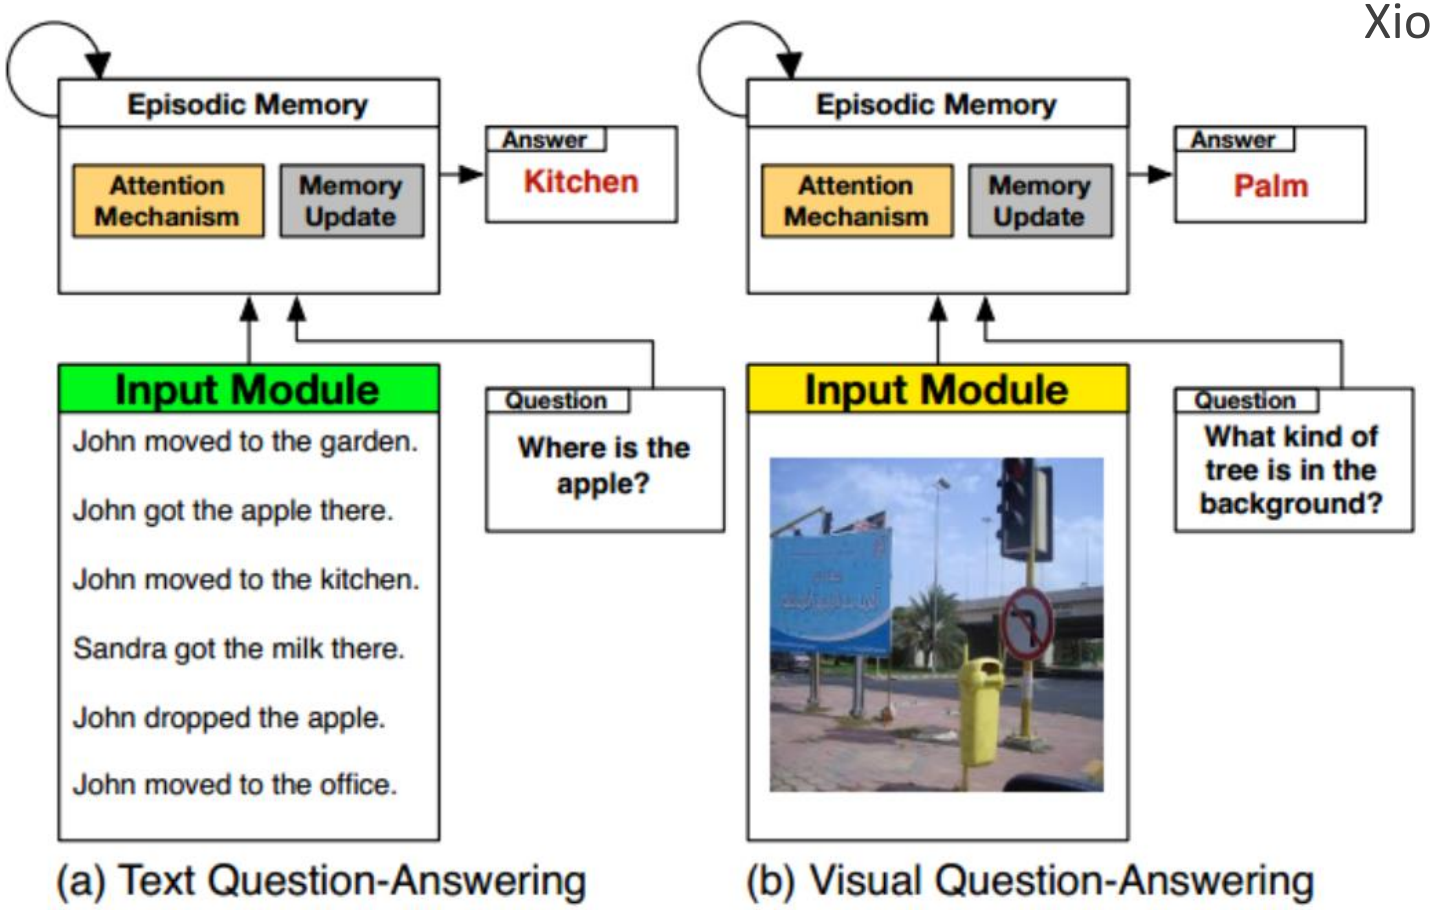

Dynamic Memory Networks (DMN)

Visual Input Module

Xiong et al. (2016)

$$\overrightarrow{f_i} = GRU_{fwd}(f_i, \overrightarrow{f_{i-1}})$$
$$\overleftarrow{f_i} = GRU_{bwd}(f_i, \overleftarrow{f_{i+1}})$$
$$\overleftrightarrow{f_i} = \overleftarrow{f_i} + \overrightarrow{f_i}$$

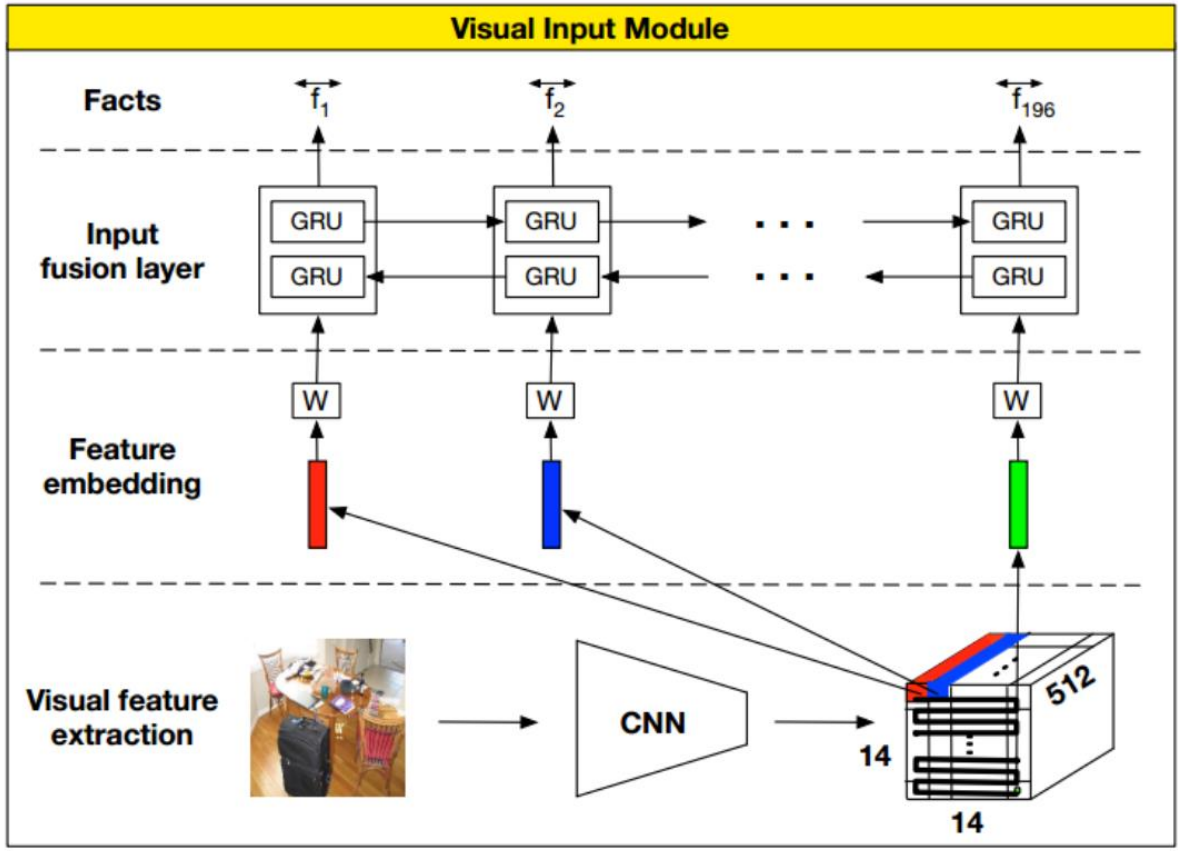

Dynamic Memory Networks (DMN)

Xiong et al. (2016)

Episodic Memory

$$z_i^t = [\overleftrightarrow{f_i} \odot q; \overleftrightarrow{f_i} \odot m^{t-1}; |\overleftrightarrow{f_i} - q|; |\overleftrightarrow{f_i} - m^{t-1}|]$$

$$Z_i^t = W^{(2)}\tanh\left(W^{(1)} z_i^t + b^{(1)}\right) + b^{(2)}$$

$$g_i^t = \text{softmax}(Z_i^t)$$

$$m^t = ReLU(W^t[m^{t-1}; c^t; q] + b)$$

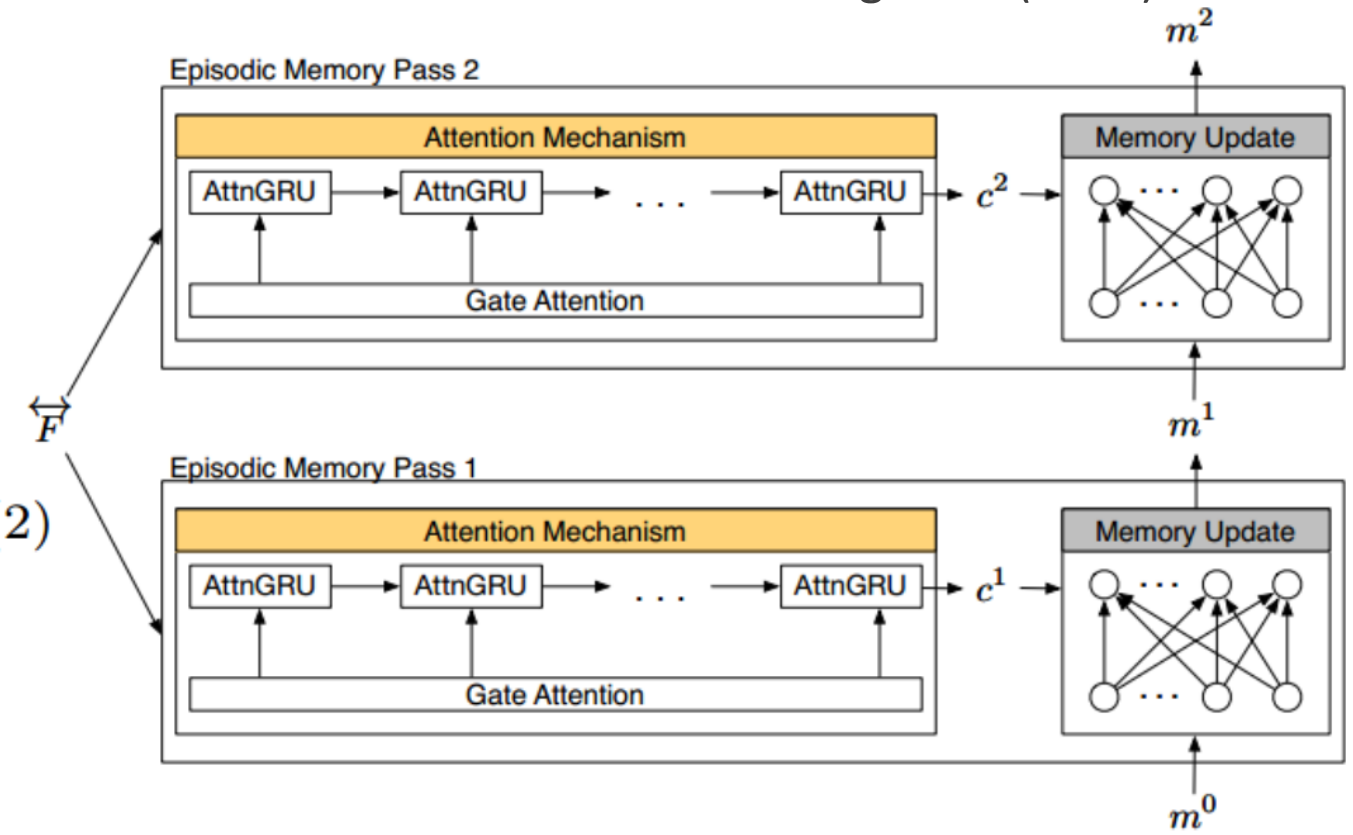

Dynamic Memory Networks (DMN)

Xiong et al. (2016)

- Image features: VGGNet (512)
- Question features: GRU (512)
- Multimodality: EWM + concatenation
- Attention mechanism: Episodic memory
- Test-dev accuracy: **60.3%**
- Test-std accuracy: **60.4%**

Multimodal Compact Bilinear (MCB)

Fukui et al. (2016)

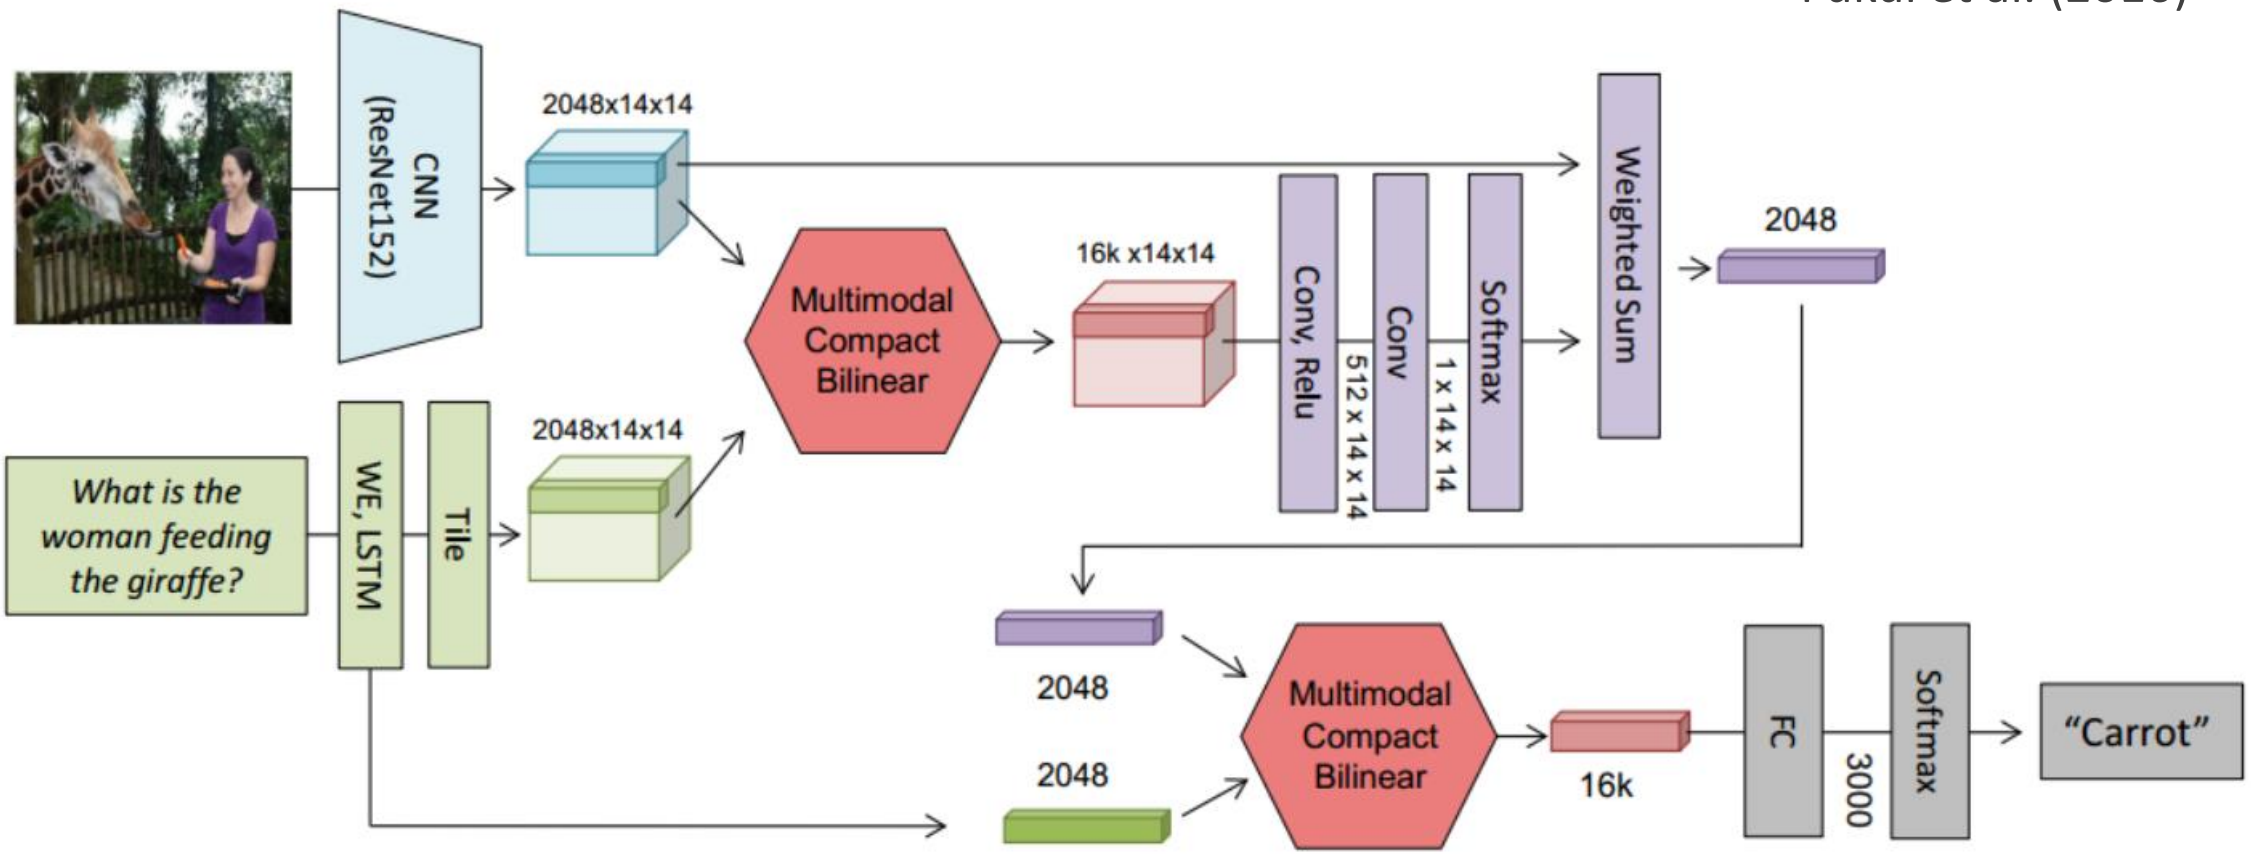

Multimodal Compact Bilinear (MCB)

Fukui et al. (2016)

Algorithm 1 Multimodal Compact Bilinear

input: $v_1 \in \mathbb{R}^{n_1}, v_2 \in \mathbb{R}^{n_2}$
output: $\Phi(v_1, v_2) \in \mathbb{R}^d$
procedure MCB(v_1, v_2, n_1, n_2, d)
    for $k \leftarrow 1 \ldots 2$ **do**
        if h_k, s_k not initialized **then**
            for $i \leftarrow 1 \ldots n_k$ **do**
                sample $h_k[i]$ from $\{1, \ldots, d\}$
                sample $s_k[i]$ from $\{-1, 1\}$
        $v'_k = \Psi(v_k, h_k, s_k, n_k)$
    $\Phi = \text{FFT}^{-1}(\text{FFT}(v'_1) \odot \text{FFT}(v'_2))$
    return Φ
procedure $\Psi(v, h, s, n)$
    $y = [0, \ldots, 0]$
    for $i \leftarrow 1 \ldots n$ **do**
        $y[h[i]] = y[h[i]] + s[i] \cdot v[i]$
    return y

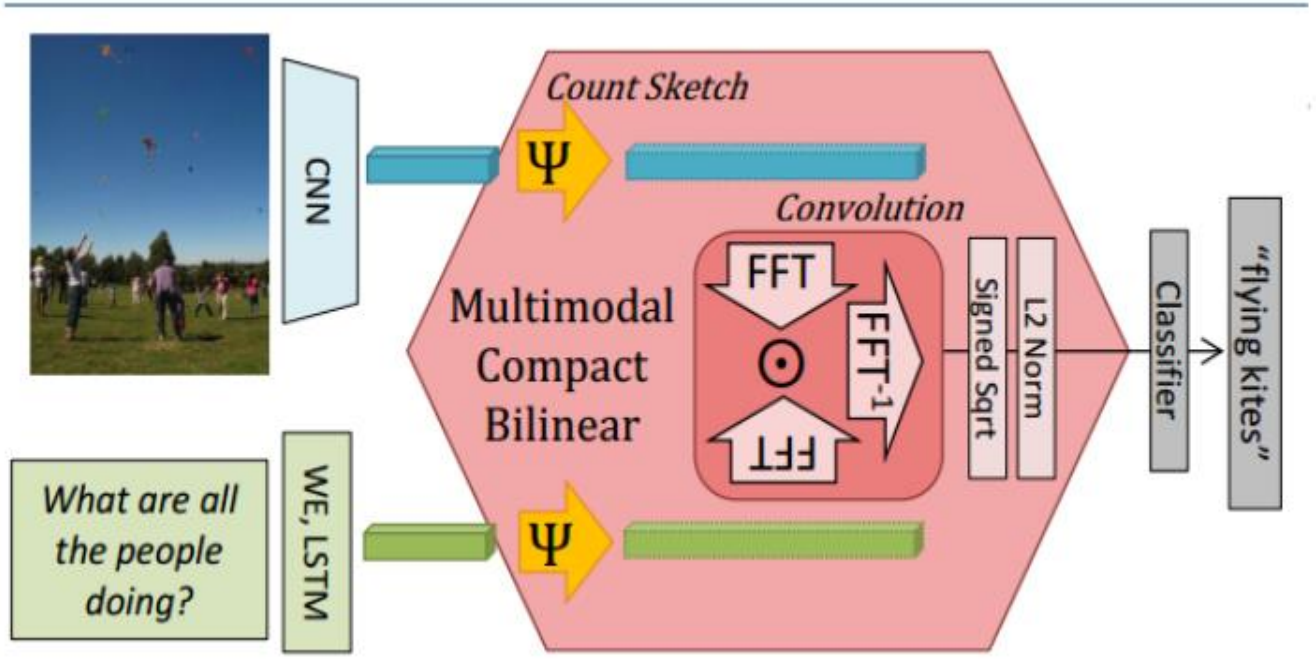

Multimodal Compact Bilinear (MCB)

Fukui et al. (2016)

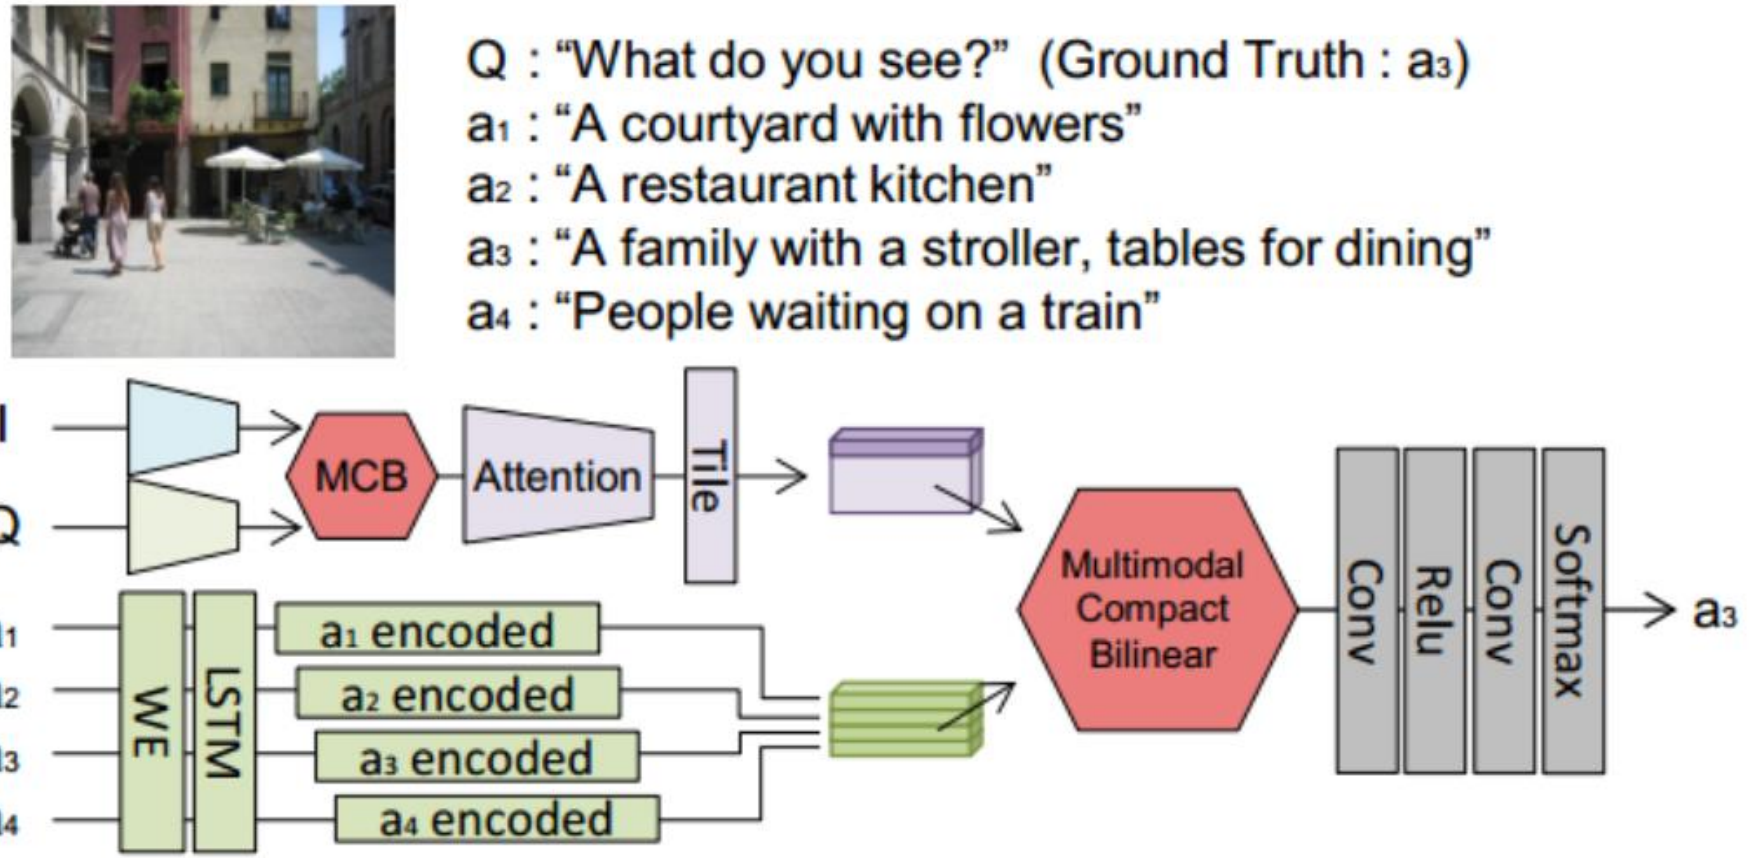

Multimodal Compact Bilinear (MCB)

Fukui et al. (2016)

- Image features: ResNet (2048)
- Question features: LSTM (1024)
- Multimodality: MCB Pooling (Count Sketch + FFT)
- Attention mechanism: MCB Pooling
- Test-dev accuracy: **66.7%**
- Test-std accuracy: **66.5%**

DualNet

Saito et al. (2016)

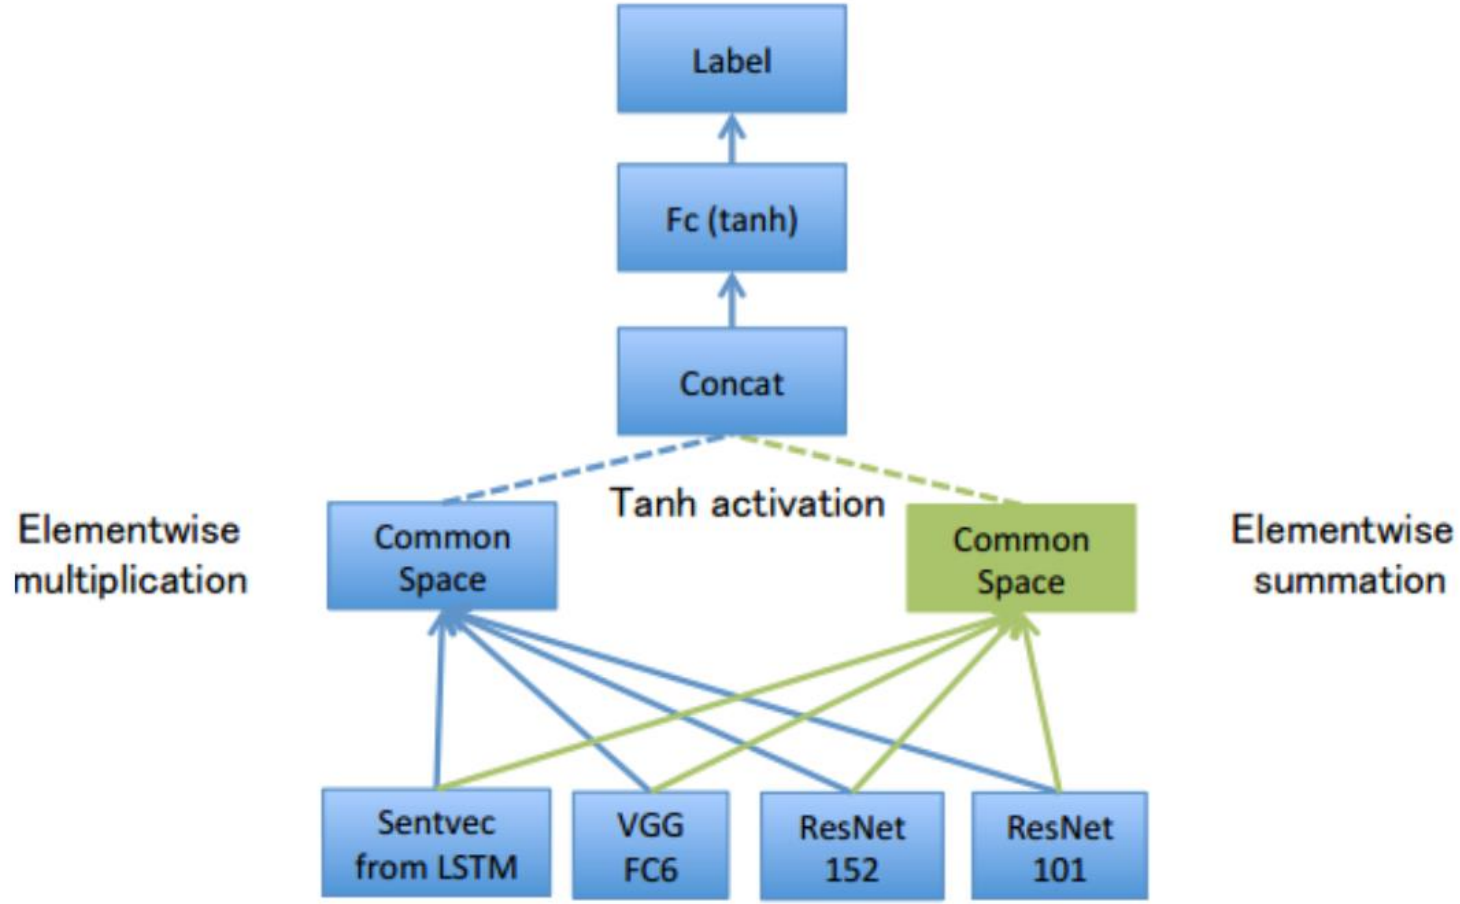

DualNet

Saito et al. (2016)

$$I_{M_1} = \tanh(W_{M_1} I_1)$$
$$I_{M_2} = \tanh(W_{M_2} I_2)$$
$$I_{M_3} = \tanh(W_{M_3} I_3)$$
$$Q_M = \tanh(W_{M_q} Q)$$
$$F_M = I_{M_1} \odot I_{M_2} \odot I_{M_3} \odot Q_M$$

$$I_{S_1} = \tanh(W_{S_1} I_1)$$
$$I_{S_2} = \tanh(W_{S_2} I_2)$$
$$I_{S_3} = \tanh(W_{S_3} I_3)$$
$$Q_S = \tanh(W_{S_q} Q)$$
$$F_S = I_{S_1} + I_{S_2} + I_{S_3} + Q_S$$

$$F = [F_M, F_S]$$
$$Output = W_{f_2} \tanh(W_{f_1} F)$$

DualNet

Saito et al. (2016)

- Image features: ResNet152 + ResNet101 + VGGNet19
- Question features: LSTM (1024)
- Multimodality: EWM + Summation
- Attention mechanism: None
- Test-dev accuracy: **61.5%**
- Test-std accuracy: **61.7%**

Algorithms Comparison

Model	Image	Question	Attention	Multimodality
LSTM Q+I [1]	VGGNet (4096)	LSTM (1024)	—	EWM
HieCoAtt [2]	VGGNet (4096)	LSTM (1024)	Image + Question	EWM
SAN [3]	VGGNet (512)	CNN (1280)	Question	EWM
MRN [4]	VGGNet (4096) ResNet (2048)	GRU (2400)	Implicit	EWM
FDA [5]	ResNet (2048)	LSTM (512)	Object Detection	EWM
DMN [6]	VGGNet (512)	GRU (512)	Episodic Memory	EWM + Concatenation
MCB [7]	ResNet (2048)	LSTM (2048)	MCB Pooling	MCB Pooling (Ψ + FFT)
DualNet [8]	ResNet152 ResNet101 VGGNet19	LSTM (1024)	—	EWM + Summation

Results Comparison

	Open-Ended					Multiple-Choice				
	test-dev				test-std	test-dev				test-std
Model	Y/N	Num	Other	All	All	Y/N	Num	Other	All	All
Human	-	-	-	-	83.3	-	-	-	-	-
LSTM Q+I [1]	78.9	35.2	36.4	53.7	54.1	79.0	35.8	43.4	57.2	-
HieCoAtt [2]	79.7	38.7	51.7	61.8	62.1	79.7	40.0	59.8	65.8	66.1
SAN [3]	79.3	36.6	46.1	58.7	58.9	-	-	-	-	-
MRN [4]	82.5	38.3	46.8	60.5	61.8	-	-	-	-	66.3
FDA [5]	81.1	36.2	45.8	59.2	59.5	81.5	39.0	54.7	64.0	64.2
DMN [6]	80.5	36.8	48.3	60.3	60.4	-	-	-	-	-
MCB [7]	**83.4**	**39.8**	**58.5**	**66.7**	**66.5**	-	-	-	**70.2**	**70.1**
DualNet [8]	82.0	37.9	49.2	61.5	61.7	**82.1**	**39.8**	**59.5**	66.7	66.7

VQA 2.0

Goyal et al. (2016)

VQA 2.0

Goyal et al. (2016)

Model	UU	UB	$B_{half}B$	BB
Prior	27.38	24.04	24.04	24.04
Language-only	48.21	41.40	41.47	43.01
LSTM Q+I [1]	54.40	47.56	49.23	51.62
HieCoAtt [2]	57.09	50.31	51.88	54.57
MCB [7]	60.36	54.22	56.08	59.14

VQA 2.0

Goyal et al. (2016)

Q: Which way is its head turned?
A: left

Q: What color is the plate?
A: blue






Future Work

- Counting mechanisms
- Apply MCB to other models
- Kernel methods
- Ensemble methods
- Emphasize image understanding (VQA 2.0)

References

[1] S. Antol, A. Agrawal, J. Lu, M. Mitchell, D. Batra, C. Lawrence Zitnick, and D. Parikh, "VQA: Visual question answering," in *The IEEE International Conference on Computer Vision (ICCV)*, December 2015.

[2] J. Lu, J. Yang, D. Batra, and D. Parikh, "Hierarchical question-image co-attention for visual question answering," in *Advances in Neural Information Processing Systems 29* (D. D. Lee, M. Sugiyama, U. V. Luxburg, I. Guyon, and R. Garnett, eds.), pp. 289–297, Curran Associates, Inc., 2016.

[3] Z. Yang, X. He, J. Gao, L. Deng, and A. J. Smola, "Stacked attention networks for image question answering," *CoRR*, vol. abs/1511.02274, 2015.

References

[4] J. Kim, S. Lee, D. Kwak, M. Heo, J. Kim, J. Ha, and B. Zhang, "Multimodal residual learning for visual QA," *CoRR*, vol. abs/1606.01455, 2016.

[5] I. Ilievski, S. Yan, and J. Feng, "A focused dynamic attention model for visual question answering," *CoRR*, vol. abs/1604.01485, 2016.

[6] C. Xiong, S. Merity, and R. Socher, "Dynamic memory networks for visual and textual question answering," *CoRR*, vol. abs/1603.01417, 2016.

[7] A. Fukui, D. H. Park, D. Yang, A. Rohrbach, T. Darrell, and M. Rohrbach, "Multimodal compact bilinear pooling for visual question answering and visual grounding," *CoRR*, vol. abs/1606.01847, 2016.

[8] K. Saito, A. Shin, Y. Ushiku, and T. Harada, "Dualnet: Domain-invariant network for visual question answering," *CoRR*, vol. abs/1606.06108, 2016.

[9] Y. Goyal, T. Khot, D. Summers-Stay, D. Batra, and D. Parikh, "Making the V in VQA matter: Elevating the role of image understanding in Visual Question Answering," in *Conference on Computer Vision and Pattern Recognition (CVPR)*, 2017.

Thank You!